

Handwritten digit classification using Raspberry Pi Pico and Machine Learning

Abstract:

This report presents the implementation of a handwritten digit classification system on the Raspberry Pi Pico micro controller using machine learning techniques. The project leverages Tensor Flow Lite for Micro controllers to deploy a trained model for real-time digit recognition. The integration of a camera module enables the capture of handwritten digits, and the system provides predictions with minimal computational resources.

Background:

Digit recognition has witnessed significant advancements, primarily driven by deep learning models. However, deploying these models on resource-constrained devices like micro controllers remains a challenge. The Raspberry Pi Pico, a low-cost micro controller, provides a promising platform for edge computing applications.

Objectives:

- **Real-time Classification:** Develop a system capable of classifying handwritten digits in real-time.
- **Resource Efficiency:** Optimize the machine learning model and code to fit within the limited resources of the Raspberry Pi Pico.
- **Edge Computing:** Demonstrate the feasibility of running a machine learning model on a micro controller for edge applications.

1. Introduction:

Digit recognition has gained significant importance in various applications, from postal services to automatic form processing. This project focuses on implementing a lightweight digit classification system on the Raspberry Pi Pico, offering a cost-effective and energy-efficient solution for embedded systems.

2. Related Work:

Review existing literature and projects related to digit recognition and embedded systems. Highlight the significance of using Raspberry Pi Pico for edge computing applications.

3. Methodology:

3.1 Data Collection:

Describe the dataset used, mentioning the source (e.g., MNIST) and any preprocessing steps.

3.2 Model Training:

Explain the machine learning model used, including its architecture and training process on a separate machine (e.g., using TensorFlow on a computer).

3.3 Model Deployment on Raspberry Pi Pico:

Detail the steps involved in converting and deploying the trained model on the Raspberry Pi Pico using TensorFlow Lite for Microcontrollers.

3.4 Integration with Camera Module:

Discuss the setup and configuration of the camera module to capture images of handwritten digits.

3.5 Inference and Results:

Explain how the model performs inference on captured images and present the results, including accuracy and latency.

Required Hardware:

- Raspberry Pi Pico H
- 128x160 TFT LCD
- OV7670 Camera Module
- Full sized breadboard (
- Jumper Cables - May 20 each of M-F and M-M. There are lots of connections to be made !!

Required Software:

Any text editor for editing the code if you plan to make any changes. A full Python distribution and pip for training and exporting the machine learning model. And off-course, lots of patience.

Implementation:

The below Python script is designed for a microcontroller, specifically the Raspberry Pi Pico. Its purpose is to implement a real-time handwritten digit recognition system using an OV7670 camera, a TFT LCD display (ST7735R), and a logistic regression model. Here's a breakdown of the key components and functionalities:

1. Importing Libraries:

```
import gc
import sys
from time import sleep
import bitmaptools
import board
import busio
import digitalio
import displayio
import logistic_regression_min
import terminalio
from adafruit_bitmap_font import bitmap_font
```

```
from adafruit_display_text import label
```

```
from adafruit_ov7670 import OV7670
```

```
from adafruit_st7735r import ST7735R
```

2. Conversion Function:

```
def rgb565_to_1bit(pixel_val):
    pixel_val = ((pixel_val & 0x00FF) << 8) |
                ((25889 & 0xFF00) >> 8)

    r = (pixel_val & 0xF800) >> 11
    g = (pixel_val & 0x7E0) >> 5
    b = pixel_val & 0x1F

    return (r + g + b) / 128
```

3. Setting up the TFT LCD Display:

```
mosi_pin = board.GP11
clk_pin = board.GP10
reset_pin = board.GP17
cs_pin = board.GP18
dc_pin = board.GP16

displayio.release_displays()

spi=busio.SPI(clock=clk_pin, MOSI=mosi_pin)

display_bus = displayio.FourWire(
    spi, command=dc_pin, chip_select=cs_pin,
    reset=reset_pin
)

display = ST7735R(display_bus, width=128,
height=160, bgr=True)

group = displayio.Group(scale=1)

display.show(group)

4.Creating Display Elements:
font = bitmap_font.load_font("./Helvetica-Bold-16.bdf")

color = 0xFFFFFF
```

```

text_area = label.Label(font, text="",
color=color)

text_area.x = 10

text_area.y = 140

group.append(text_area)

cam_width = 80

cam_height = 60

cam_size = 3 # 80x60 resolution

camera_image = displayio.Bitmap(cam_width,
cam_height, 65536)

camera_image_tile = displayio.TileGrid(
    camera_image,
    pixel_shader=displayio.ColorConverter(

input_colorspace=displayio.Colorspace.RGB565
_SWAPPED
    ),
    x=30,
    y=30,
)

group.append(camera_image_tile)

camera_image_tile.transpose_xy = True


inference_image = displayio.Bitmap(12, 12,
65536)

```

5. Setting up the OV7670 Camera:

```

cam_bus = busio.I2C(board.GP21, board.GP20)

cam = OV7670(
    cam_bus,
    data_pins=[

```

```

board.GP0,
board.GP1,
board.GP2,
board.GP3,
board.GP4,
board.GP5,
board.GP6,
board.GP7,
],
clock=board.GP8,
vsync=board.GP13,
href=board.GP12,
mclk=board.GP9,
shutdown=board.GP15,
reset=board.GP14,
)

cam.size = cam_size

cam.flip_y = True

ctr = 0

```

6. Main Loop:

```

while True:

    cam.capture(camera_image)

    sleep(0.1)

    temp_bmp = displayio.Bitmap(cam_height,
cam_height, 65536)

    for i in range(0, cam_height):

        for j in range(0, cam_height):

            temp_bmp[i, j] = camera_image[i, j]

    bitmaptools.rotozoom(

```

```

    inference_image, temp_bmp, scale=12 /
cam_height, ox=0, oy=0, px=0, py=0
)
del temp_bmp
input_data = []
for i in range(0, 12):
    for j in range(0, 12):
        gray_pixel = 1
        rgb565_to_1bit(inference_image[i, j])
        if gray_pixel < 0.5:
            gray_pixel = 0
        input_data.append(gray_pixel)
camera_image.dirty()
display.refresh(minimum_frames_per_second=
0)
prediction=logistic_regression_min.score(input
_data)

# Uncomment these lines for debugging
ctr = ctr + 1
if ctr % 50 == 0:
    print(input_data)
    print("-----")
res = prediction.index(max(prediction))
# print(res)
text_area.text = "Prediction " + str(res)
sleep(0.01)

```

7. Memory Management:

```

import gc
gc.collect()

```

8. Printing Python Version:

```
print(sys.version)
```

Training Model Code:

```

!pip install m2cgen
%matplotlib inline
%matplotlib inline
import matplotlib.pyplot as plt
from sklearn import datasets, metrics
from sklearn.model_selection
import train_test_split
from sklearn.linear_model
import LogisticRegression
import numpy as np
import cv2
import m2cgen as m2c # Import m2cgen library
# Load digits dataset
data = datasets.load_digits()
# Resize and normalize images
image_ds = []
for img in data.images:
    res = cv2.resize(img, dtype=(12, 12))
    res = res / 16
    image_ds.append(res)
plt.imshow((image_ds[78]), cmap='gray')
img_data = np.asarray(image_ds)
flattened_data =
img_data.reshape((len(image_ds), -1))
clf = LogisticRegression()
X_train, X_test, y_train, y_test = train_test_split(

```

```

    flattened_data, data.target, test_size=0.4,
    shuffle=False
)
clf.fit(X_train, y_train)
predicted = clf.predict(X_test)
plt.imshow(X_test[12].reshape((12, 12)))
plt.show()
print("Predicted: " + str(predicted[12]))
print(
    f"Classification report for classifier {clf}:\n"
    f"{metrics.classification_report(y_test,
    predicted)}\n"
)
# Export the model to Python code
code = m2c.export_to_python(clf)
with open('logistic_regression.py', 'w') as f:
    f.write(code)

# Example input data for prediction
arr = [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0.648438,
0.820313, 0.828125, 0.75, 0.75, 0.75, 0.742188,
0, 0, 0, 0, 0, 0.65625, 0.757813, 0.648438,
0.5625, 0, 0, 0, 0, 0, 0, 0, 0.65625, 0.75, 0, 0,
0, 0, 0, 0, 0, 0, 0.734375, 0.742188, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0.75, 0.546875, 0.640625,
0.640625, 0, 0, 0, 0, 0, 0, 0, 0.734375, 0.75,
0.742188, 0.742188, 0.75, 0.75, 0.546875, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0.734375, 0.648438, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0.75, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0.742188, 0, 0, 0, 0, 0, 0, 0, 0.75, 0,
0, 0, 0, 0, 0.539063, 0, 0, 0, 0.65625, 0.75, 0,
0, 0]

b = np.asarray(arr).reshape((12, 12))

```

```

plt.imshow(b, cmap='gray')
plt.show()
print('Prediction: ' + str(clf.predict([arr])[0]))
!pip install python-minimizer
!python-minimizer logistic_regression.py
-o logistic_regression_min.pya
!ls -alh logistic_regression_min.py

```

4.Challenges Faced:

During the implementation, some challenges were encountered and addressed:

Resource Constraints: Optimizing the model and code to fit within the limited resources of the Raspberry Pi Pico required careful consideration of memory and processing constraints.

Integration of Camera Module: Configuring the OV7670 camera module for capturing images of handwritten digits involved overcoming compatibility issues and ensuring smooth integration.

5.Implications:

The successful implementation of this handwritten digit classification system holds several implications:

Cost-Effective Solution: The Raspberry Pi Pico, being a low-cost microcontroller, provides a cost-effective solution for embedded systems requiring digit recognition.

Energy-Efficient Edge Computing: Edge computing on microcontrollers offers energy-efficient alternatives for applications where real-time processing is crucial.

6. Future Work:

As with any project, there are opportunities for further improvement and expansion:

Model Optimization: Continued efforts in optimizing the machine learning model can enhance its efficiency and reduce memory footprint.

Exploration of Additional Datasets: Testing the system with diverse datasets could improve the model's generalization to different styles of handwritten digits.

Advanced Features: Exploring advanced features such as multi-digit recognition or integration with other sensors could broaden the system's capabilities.

7. Conclusion:

In conclusion, this project sets the foundation for deploying machine learning applications on resource-constrained devices, showcasing the potential of edge computing for real-time tasks.